



School of Computer Science

COMP47470

Lab 1
CLI (Bash) for Big Data

Teaching Assistant:	Patrick Cormac English
Coordinator:	Dr Anthony Ventresque
Date:	Tuesday 21 st September, 2021
Total Number of Pages:	9

Linux is an Operating System that is widely used in industry - in particular the vast majority of server-side systems (web servers, app servers etc.) run on Linux machines. In this lab, you will be using a container (a lightweight virtual machine) running Linux on your own system. You will then play with some basic Linux system commands.

1 Setting our Linux Container Up

In this module we will work with our Big Data technologies in *containers* (using Docker), i.e., applications bundled with all their environment, libraries and dependencies. The containers are portable and independent from Operating Systems - and much lighter than VMs. Note that containers don't contain a full OS like VMs do and as such do not provide as many functionalities - but they are fine when you're only looking for a solution that runs applications in a simple software container that can be shipped to various environments.

1.1 Installing docker

Get docker for your operating system *here* and follow the instructions to install on your system. **NB - Docker for Windows will likely require the installation of the Windows Linux Subsystem, which is semi-convoluted. If you are a Windows user, we would ask that you have some patience in this regard as there are several points of failure in the installation process - the TA is experienced with Windows installs and will ensure that we sort something out for you!**

1.2 Our first container

Now that docker is installed on our system, we want to start our first container. We will use a simple image that sets up an Ubuntu container so we have a Linux environment to play in. Later, we will use more complex images that contain the different tools we will be needing. First let's create and start our container. In a terminal (cmd/powershell in Windows, the terminal app in mac/linux), type (the \$ only shows the beginning of the line, do not type it):

```
$ docker run -tid --name aNameYouLike registry.gitlab.com/roddhjav/ucd-bigdata/bash
```

We use the `docker` command, and tell docker to **run** the `ubuntu` image and give our container the name "aNameYouLike" (you can change this, but remember the name you chose). Let's check that our container is running:

```
$ docker ps
```

This should show you the container you just created, running the `ubuntu` docker image, with some information such as its ID or how long it has been running. We can start and stop our container so it only runs when we need it:

```
$ docker stop aNameYouLike
$ docker start aNameYouLike
```

When the container is running (just after we did **run** or after doing **start**), we can then run `bash` inside the container:

```
$ docker exec -it aNameYouLike bash
```

This will open a bash command line in the container where you can do the rest of the practical.

2 Using man to Get Help

Nearly every command and application in Linux will have a man (manual) page, so finding them is as simple as typing `man command` to bring up a manual entry for the specified command.

- For example, `man mv` will bring up the `mv` (move) manual.
- Move up and down the man file with the arrow keys, and get back to the command prompt with "q".
- `man man` will bring up the manual entry for the `man` command, which is a good place to start!
- `man intro` is especially useful - it displays the "Introduction to user commands" which is a well-written, fairly brief introduction to the Linux command line.
- Some software developers prefer `info` to `man` (for instance, GNU developers), so if you find a very widely used command or app that doesn't have a man page, it's worth checking for an info page (Please note: info is not installed on your container).
- Virtually all commands understand the `-h` (or `--help`) option which will produce a short usage description of the command and its options, then exit back to the command prompt. Try `man -h` or `man --help` to see this in action.
- man pages can be lengthy, so if you are looking for a specific option etc. it could be useful to look up some word using the syntax `/word` – and then hit the key `n` to move to the next occurrence.
- If you aren't sure which command or application you need to use, you can try searching the manual pages. Each manual page has a name and a short description.
 - If you know part of the command name, use the following command: `whatis -r <string>`. For example try with the following: `whatis -r cpy`
 - To search the names or descriptions for `<string>` enter: `apropos -r <string>`. For example, `apropos -r "copy files"` will list manual pages whose names or descriptions contain copy files.

3 File & Directory Commands

The Linux hierarchy is typical of Unix systems (with some variations depending on the specific distributions). For the moment you just need to know that the file system is a tree that starts at the root (represented with the symbol `/`). Note that if you are familiar with DOS/Windows the path delimiter is the forward slash and not the backward slash... A path then looks like this in Linux: `/var/log/auth.log`. This leads to the file `auth.log` in the folder `log` in the folder `var` which is right after the root of the file system.

- The tilde (~) symbol stands for your home directory. If you are anthony, then the tilde (~) stands for `/home/anthony`. So `/home/anthony/myFile` and `~/myFile` point to the same file.
- **pwd**: The `pwd` command will allow you to know in which directory you're located (`pwd` stands for "print working directory"). Example: `pwd` in the Desktop directory will show `~/Desktop`.
- **ls**: The `ls` command will show you ('list') the files in your current directory. Used with certain options, you can see the size of files, when files were created, and permissions for files. Example: `ls ~` will show you the files that are in your home directory. Check some of the options of `ls`.
- **cd**: The `cd` command will allow you to change directories. When you open a terminal you will be in your home directory. To move around the file system you will use `cd`. Examples:
 - To navigate into the root directory, use `cd /`
 - To navigate to your home directory, use `cd` or `cd ~`
 - To navigate up one directory level, use `cd ..`
 - To navigate to the previous directory (or back), use `cd -`
 - To navigate through multiple levels of directory at once, specify the full directory path that you want to go to. For example, use, `cd /var/log` to go directly to the `/log` subdirectory of `/var/`. As another example, `cd ~/Desktop` will move you to the Desktop subdirectory inside your home directory.
- **cp**: The `cp` command will make a copy of a file for you. Example: `cp file foo` will make an exact copy of "file" and name it "foo", but the file "file" will still be there. If you are copying a directory, you must use `cp -r directory foo` (copy recursively). (To understand what "recursively" means, think of it this way: to copy the directory and all its files and subdirectories and all their files and subdirectories of the subdirectories and all their files, and on and on, "recursively")
- **mv**: The `mv` command will move a file to a different location or will rename a file. Examples are as follows: `mv file foo` will rename the file "file" to "foo". `mv foo ~/Desktop` will move the file "foo" to your Desktop directory, but it will not rename it. You must specify a new file name to rename a file.
- **rm**: Use this command to remove or delete a file in your directory.
- **rmdir**: The `rmdir` command will delete an empty directory. To delete a directory and all of its contents recursively, use `rm -r` instead.
- **mkdir**: The `mkdir` command will allow you to create directories. Example: `mkdir music` will create a directory called "music".

Let's start with some exercises!

Exercise

- Explore the filesystem tree using `cd`, `ls` and `pwd`. Look in `/bin`, `/usr/bin`, and `/sbin`. What do you see?
- Make a directory called `BigData` (`mkdir`) in your home directory and then rename it to `COMP47470` (`mv`)
- Make a sub-directory `lab1` and in it start a text file `lab1_commands.txt` using **nano** (command `nano fileName` where you can write all the examples and exercises you'll do today (or anything you like). The controls for nano are displayed at the bottom of the terminal, `^` means the Ctrl key, so for example `^X` means press Ctrl and x.
- Copy the file `/etc/passwd` into your home directory (`cp`).
- use the following command `ls -l /var/log/ | wc -l` to get the number of files in `/var/log/`. How does `wc` work? (check this in the `man` page corresponding to `wc`). We will spend some time on the `|` symbol next week – but do you have an idea of what it does? try `ls -l /var/log/` alone to give you an idea of what happens.

4 Examining files. *Why use more if you have less*

"cat" stands for conCATenate. You can use this command to dump an entire text file to the screen.

```
$ cat /etc/passwd
```

If the text file is too long, you might find that it scrolls past too quickly and you cannot see the beginning of the file anymore. In which case, you can use either the "more" or "less" command. For example

```
$ more /etc/passwd
```

Or

```
$ less /etc/passwd
```

Both these commands perform similar functions as they allow to see the text file one page at a time. You use the spacebar to continue paging, `<enter>` key will move down one line, and `"q"` to quit.

"less" has actually more features than "more" :) The most useful feature is that it can scroll backwards (or up) whereas "more" cannot. Press `"h"` (while in the program) to see more options. Another interesting option is the search option, which is similar to the one you've seen when we presented `man`. If you are looking for a specific string in a text file use the syntax `/string -` and then hit the key `n` to move to the next occurrence.

5 Showing part of a file. *Can you make heads or tails of it?*

Note: `/etc/passwd` is a file storing some login info for a user: user number, group number, shell at login, home directory etc. It also contains OS processes considered as "users". Use commands `whoami`, `id`, `groups` to try to get some of the information contained in `/etc/passwd`.

`head` and `tail` are two opposite commands, showing the beginning or the end of a file respectively. Try the following commands – what do they give you?

```
$ head /etc/passwd
```

```
$ tail /etc/passwd
```

Both commands have various options that can be powerful – and a little complicated:

- `$ head -n 3 /etc/passwd` shows the first 3 lines
- `$ tail -n 5 /etc/passwd` shows the last 5 lines
- `$ tail -n +4 /etc/passwd` shows from the fourth line to the end
- `$ head -n -1 /etc/passwd` shows from the second line to the tenth line

A classic use of `tail` consists in showing the content of a dynamic file, that is evolving over time. `/var/log/` contains a lot of files that store log messages, i.e., information about what's happening in the system. Try `tail -f /var/log/syslog`.

To get out of the `tail` program press the combination of keys *control* + *c*

Exercise

- Combine `head` and `tail` using a pipe to print 15th line to 20th line in `/etc/passwd`.
- Combine `head` and `tail` using a pipe to print the 7th line in `/etc/passwd`

Echo and Special Characters.

The goal of this section is to play with special characters and to show how they impact the execution of commands, and can be disabled conveniently. If you need to interrupt/stop the execution of a command (if you're stuck), you can use the combination of keys *control* + *c*.

- print on the screen a simple word, for instance "hello" (use `echo`).
- `echo` can print several strings - print "hello everybody"
- you can also use `echo` to print non alphabetical characters on screen (as long as they do not belong to the list of special characters). Use `echo` to print "Hello to the 2 of you!"
- Disable the special character `#` in the following string (using the character `\`) and print it on screen using `echo`: "`#` is a very useful character in bash"

- Likewise for the following: "# is more useful than \ in bash"
- Print the following using echo: "# is more useful than \ in bash but less than *"
- When you need to disable a lot of special characters, you can use simple and double quotes to disable all (simple quotes) or some (double quotes - do not disable \$, ' and \) of them. Print the following two sentences using simple then double quotes: "The \$PATH variable is very important."

Selecting Columns. *Charles I of England?*

Linux command `cut` is used for text processing. You can use this command to extract portions of text from a file by selecting columns.

Option `-cN` extracts only column (character) `N` of each line from a file:

```
$ cut -c2 /etc/passwd
```

will extract only the second character of each line (the 2nd column of each line).

Range of characters can also be extracted from a file by specifying start and end position delimited with `-`.

```
$ cut -c2-5 /etc/passwd
```

Either start position or end position can be passed to the `cut` command with the `-c` option.

The following command specifies only the start position before the `'-'`. This example extracts from the 10th character to the end of each line.

```
$ cut -c10- /etc/passwd
```

The following command specifies only the end position after the `'-'`. This example extracts 10 characters from the beginning of each line.

```
$ cut -c-10 /etc/passwd
```

Instead of selecting `x` number of characters, if you like to extract a whole field, you can combine options `-f` and `-d`. Option `-f` specifies which field you want to extract, and option `-d` specifies the field delimiter that is used in the input file.

The following example displays only the first field of each line from the `/etc/passwd` file using the field delimiter `:` (colon). In this case, the 1st field is the username.

```
$ cut -d":" -f1 /etc/passwd
```

You can also extract more than one field from a file. The example below displays the username and home directory of users.

```
$ cut -d":" -f1,6 /etc/passwd
```

To display a range of fields specify start field and end field as shown below. In this example, we are selecting fields 1 through 4, 6 and 7 (fields 1, 2, 3, 4, 6 and 7).

```
$ cut -d":" -f1-4,6,7 /etc/passwd
```

Sorting. *The Last Shall Be First*

The sort command rearranges the lines in a text file so that they are sorted lexicographically.

```
$ sort /etc/passwd
```

These are the default sorting rules:

- a number is before a letter.
- letters follow their order in the alphabet.
- a lowercase letter is before a same uppercase letter.

You can change the default sorting rules by providing a suitable parameter. E.g..

- Reverse sort:

```
$ sort -r /etc/passwd
```

- Ignore case:

```
$ sort -f /etc/passwd
```

You can also concatenate and sort several files at once:

```
$ sort /etc/passwd /etc/group
```

You can also save the result of the sort in a another file:

```
$ sort /etc/passwd -o result.txt  
or  
$ sort /etc/passwd > result.txt  
or  
$ sort /etc/passwd >> result.txt
```

The -o option is not available for all commands. Using > and >> works for any command. > will overwrite result.txt if it already exists while >> will append the output to result.txt if it already exists. Both create the file if it does not exist.

Duplicates. *Castor Troy: It's like looking in a mirror, only not.*

This command filters duplicate adjacent lines.

```
$ uniq /etc/passwd
```

It can filter them even by ignoring the case:

```
$ uniq -i /etc/passwd
```

You can report the duplicate lines by:

```
$ uniq -d /etc/passwd
```

You can also print the number of occurrences of each line:

```
$ uniq -c /etc/passwd
```


Grep. *Look it up!*

Grep prints out the lines containing a certain string.

```
$ grep root /etc/passwd
```

Options:

- -c: only gives the number of matching lines
- -v: Shows only the lines that do not match the pattern. Inverted search.
- -i: ignore case
- -n: gives the line number as well as the matching lines.

Getting files on and from the container

To get files out of and into your container you can use `docker cp` (documentation [here](#)). `docker cp` works like `cp` but lets us copy file to/from our machine (the host) from/to our container. For example, if we want to copy the file `myFile` that is in the `/root` folder on the `myContainer` container to the current directory on the host, we can do:

```
$ docker cp myContainer:/root/myFile ./
```

Exercise 1

- What is the output of the commands:
 - `echo {0..9}`
 - `echo 1.{0..9}`
 - `echo {A..C}{0..2}`
- Use `sort` to sort the content of `/etc/passwd`
- Use `grep` to list only the line from `/proc/meminfo` that shows total memory (the name of the variable is `MemTotal`).
- Use `grep` to find all the lines with "to" in the `syslog` file (`/var/log/syslog_public`). Then redirect this output into a file in your home directory.
- List all of the files in a directory in reverse date order, so that the most recent is at the bottom (`ls`).

Exercise 2

In this exercise, we will be exploring a dataset containing information regarding higher education institutions. Follow these steps using the commands described above.

1. Download the dataset as `unirank.csv` using `wget` (capital o option):

```
wget -O unirank.csv http://csserver.ucd.ie/~thomas/unirank.csv
```

2. Explore the file and its structure (`cat`, `head`, `tail`, `less`, ...)
3. Find all the "colleges" in the list (e.g. University *College* Dublin, `thinkgrep`).
4. List and group all the HEIs by state (`cut`, `sort`).
5. Find the number of HEIs per state (`cut`, `sort`, `uniq`). (which state has the most institutes in the dataset?)
6. Create a list (x, y) of couples containing the Tuitions and Fees (x) and the rank (y) for each institute. Get this file to your host machine and plot the values using Excel, gnuplot, etc. What do you notice?